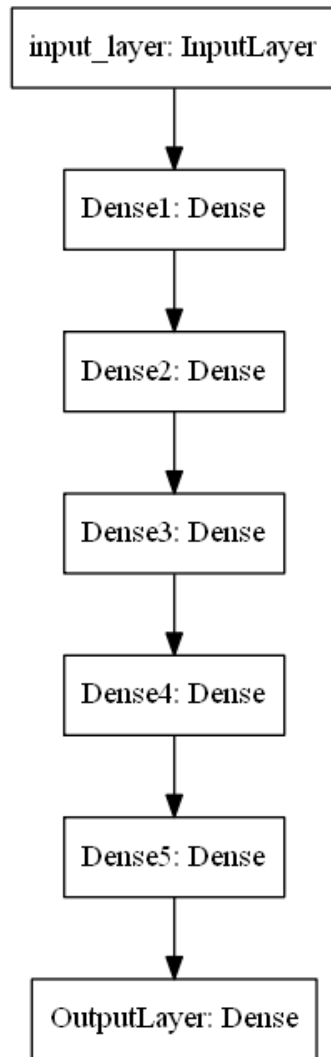


1. Download the data from [here](#)
2. Code the model to classify data like below image



3. Write your own callback function, that has to print the micro F1 score and AUC score after each epoch. **\*\*DONE\*\***
4. Save your model at every epoch if your validation accuracy is improved from previous epoch. **\*\* KINDA DONE\*\***

5. you have to decay learning based on below conditions

Cond1. If your validation accuracy at that epoch is less than previous epoch accuracy, you have to decrease the learning rate by 10%. DONE

Cond2. For every 3rd epoch, decay your learning rate by 5%.DONE

6. If you are getting any NaN values(either weights or loss) while training, you have to terminate your training. DONE

7. You have to stop the training if your validation accuracy is not increased in last 2 epochs. DONE

8. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)

9. use cross entropy as loss function

10. Try the architecture params as given below.

#### **Model-1**

1. Use tanh as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

#### **Model-2**

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use RandomUniform(0,1) as initializer.
3. Analyze your output and training process.

#### **Model-3**

1. Use relu as an activation for every layer except output layer.
2. use SGD with momentum as optimizer.
3. use he\_uniform() as initializer.
3. Analyze your output and training process.

#### Model-4

1. Try with any values to get better accuracy/f1 score.

```
%load_ext tensorboard
```

```
!rm -rf ./logs/
```

```
pip install tensorflow-addons
```

```
Collecting tensorflow-addons
  Downloading tensorflow_addons-0.14.0-cp37-cp37m-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (1.1 MB)
    |████████████████████████████████████████| 1.1 MB 5.2 MB/s
Requirement already satisfied: typeguard>=2.7 in /usr/local/lib/python3.7/dist-packages (from tensorflow-addons) (2.7.1)
Installing collected packages: tensorflow-addons
Successfully installed tensorflow-addons-0.14.0
```

```
import tensorflow as tf
import pandas as pd
from sklearn.model_selection import train_test_split
import datetime
```

```
from tensorflow.keras.layers import Dense, Input, Activation
import keras
from tensorflow.keras.models import Model
import tensorflow_addons as tfa
from keras import metrics
from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler, ReduceLROnPlateau, TerminateOnNaN, ModelCheckpoint
import random as rn
```

```
import numpy as np
from keras.callbacks import Callback
from sklearn.metrics import confusion_matrix, f1_score, precision_score, recall_score, roc_auc_score
```

```
data = pd.read_csv('data.csv')
```

```
X = data[['f1', 'f2']]
y = data['label']
```

```
print(X.shape)
print(y.shape)
```

```
(20000, 2)
(20000,)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
```

```
def changedLearningRate(epoch, lr):
    """In this function, we decrease the learning rate by 5% every third epoch"""
    if ((epoch+1)%3 == 0):
        lr = 0.95 * lr
    return float(lr)
```

```
class Metrics(Callback):
```

```
    def on_train_begin(self, logs={}):
        self.val_f1s = []
        self.val_recalls = []
        self.val_precisions = []
        self.val_aucs = []
```

```
    def on_epoch_end(self, epoch, logs={}):
        val_predict = (np.asarray(self.model.predict(X_test))).round()
        val_targ = y_test
        _val_f1 = f1_score(val_targ, val_predict)
        _val_recall = recall_score(val_targ, val_predict)
        _val_precision = precision_score(val_targ, val_predict)
        _val_auc = roc_auc_score(val_targ, val_predict)
        self.val_f1s.append(_val_f1)
        self.val_recalls.append(_val_recall)
        self.val_precisions.append(_val_precision)
```

```

self.val_aucs.append(_val_auc)

print("val_f1: %f - val_precision: %f - val_recall %f - val_auc %f" %(_val_f1, _val_precision, _val_recall, _val_auc))

return

metrics = Metrics()
#source: https://medium.com/@thongonary/how-to-compute-f1-score-for-each-epoch-in-keras-a1acd17715a2

```

## ▼ Model 1

```

input_layer = Input(shape=(2, ))

layer1 = Dense(784, activation='tanh', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(input_layer)
layer2 = Dense(512, activation='tanh', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer1)
layer3 = Dense(256, activation='tanh', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer2)
layer4 = Dense(64, activation='tanh', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer3)
layer5 = Dense(32, activation='tanh', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer4)

output = Dense(1, activation='sigmoid', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer5)

model1 = Model(inputs = input_layer, outputs = output)
optimizer = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.1, nesterov=True)

model1.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

log_dir="logs/fit/model1/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
reduce_lr = ReduceLROnPlateau(monitor = 'val_accuracy', factor=0.9, patience=1, verbose=1, mode='auto')

terminate_tr = TerminateOnNaN()
earlystop = EarlyStopping(monitor='val_accuracy', patience=2, verbose=1, mode='max')

lrschedule = LearningRateScheduler(changedLearningRate, verbose=1)
model1.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test), validation_split=0.2, callbacks=[checkpoint, reduce_lr, terminate_tr,

    WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
    Epoch 1/10

```

```
Epoch 00001: LearningRateScheduler setting learning rate to 0.10000000149011612.
3/335 [.....] - ETA: 43s - loss: 6.1710 - accuracy: 0.5521WARNING:tensorflow:Callback method `on_train_batch_begin` is
WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0048s vs `on_train_batch_end` time: 0.0
335/335 [=====] - 5s 6ms/step - loss: 0.7967 - accuracy: 0.5089 - val_loss: 0.7089 - val_accuracy: 0.5086

Epoch 00001: val_accuracy improved from -inf to 0.50858, saving model to model_save/weights-01-0.5086.hdf5
val_f1: 0.491618 - val_precision: 0.491768 - val_recall 0.491469 - val_auc 0.494529
Epoch 2/10

Epoch 00002: LearningRateScheduler setting learning rate to 0.10000000149011612.
335/335 [=====] - 1s 4ms/step - loss: 0.7062 - accuracy: 0.4989 - val_loss: 0.6951 - val_accuracy: 0.5086

Epoch 00002: val_accuracy did not improve from 0.50858

Epoch 00002: ReduceLROnPlateau reducing learning rate to 0.09000000134110452.
val_f1: 0.491618 - val_precision: 0.491768 - val_recall 0.491469 - val_auc 0.494529
Epoch 3/10

Epoch 00003: LearningRateScheduler setting learning rate to 0.08550000339746475.
335/335 [=====] - 1s 4ms/step - loss: 0.7037 - accuracy: 0.4998 - val_loss: 0.6948 - val_accuracy: 0.5086

Epoch 00003: val_accuracy did not improve from 0.50858

Epoch 00003: ReduceLROnPlateau reducing learning rate to 0.07695000171661377.
val_f1: 0.491618 - val_precision: 0.491768 - val_recall 0.491469 - val_auc 0.494529
Epoch 00003: early stopping
<keras.callbacks.History at 0x7fa011530c50>
```



```
%load_ext tensorboard
```

```
The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

```
!kill 197
```

```
/bin/bash: line 0: kill: (197) - No such process
```

```
%tensorboard --logdir logs/fit/model1/
```

☐ Show data download links

☐ Ignore outliers in chart scaling

 Tooltip sorting method: default ▼

Smoothing

Horizontal Axis

STEP

RELATIVE

WALL

Runs


Write a regex to filter runs

☐ 20211009-070754/train

☐ 20211009-070754/validation

TOGGLE ALL RUNS

logs/fit/model1/



run to download ▼

epoch\_loss



epoch\_lr

epoch\_lr  
tag: epoch\_lr

run to download ▼

## ▼ Model 2

```
input_layer = Input(shape=(2, ))
```

```
layer1 = Dense(784, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(input_layer)
```

```
layer2 = Dense(512, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer1)
```

```
layer3 = Dense(256, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer2)
```

```
layer4 = Dense(64, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer3)
```

```
layer5 = Dense(32, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer4)
```

```
layer5 = Dense(32, activation='relu', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer4)
```

```
output = Dense(1, activation='sigmoid', kernel_initializer=tf.keras.initializers.random_uniform(0, 1))(layer5)
```

```
model2 = Model(inputs = input_layer, outputs = output)
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.9, nesterov=True)
```

```
model2.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])
```

```
log_dir="logs/fit/model2/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
```

```
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)
```

```
filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
```

```
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
```

```
reduce_lr = ReduceLRonPlateau(monitor = 'val_accuracy', factor=0.9, patience=1, verbose=1, mode='auto')
```

```
terminate_tr = TerminateOnNaN()
```

```
earlystop = EarlyStopping(monitor='val_accuracy', patience=2, verbose=1, mode='max')
```

```
lrschedule = LearningRateScheduler(changedLearningRate, verbose=1)
```

```
model2.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test), validation_split=0.2, callbacks=[checkpoint, reduce_lr, terminate_tr,
```

```
WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.
```

```
Epoch 1/10
```

```
Epoch 00001: LearningRateScheduler setting learning rate to 0.10000000149011612.
```

```
3/335 [.....] - ETA: 37s - loss: 271526528.0000 - accuracy: 0.5000WARNING:tensorflow:Callback method `on_train_batch_b
```

```
WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0039s vs `on_train_batch_end` time: 0.0
```

```
335/335 [=====] - 2s 6ms/step - loss: 2431580.7500 - accuracy: 0.5023 - val_loss: 0.6955 - val_accuracy: 0.4966
```

```
Epoch 00001: val_accuracy improved from -inf to 0.49664, saving model to model_save/weights-01-0.4966.hdf5
```

```
val_f1: 0.000000 - val_precision: 0.000000 - val_recall 0.000000 - val_auc 0.500000
```

```
Epoch 2/10
```

```
Epoch 00002: LearningRateScheduler setting learning rate to 0.10000000149011612.
```

```
16/335 [>.....] - ETA: 1s - loss: 0.6947 - accuracy: 0.4824/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_class
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
335/335 [=====] - 2s 4ms/step - loss: 0.6949 - accuracy: 0.4949 - val_loss: 0.6958 - val_accuracy: 0.4966
```

```
Epoch 00002: val_accuracy did not improve from 0.49664
```

```
Epoch 00002: ReduceLRonPlateau reducing learning rate to 0.09000000134110452.
```

```
val_f1: 0.000000 - val_precision: 0.000000 - val_recall 0.000000 - val_auc 0.500000
```

```
Epoch 3/10
```

```
Epoch 00003: LearningRateScheduler setting learning rate to 0.08550000339746475.
```

```
15/335 [>.....] - ETA: 1s - loss: 0.6925 - accuracy: 0.5271/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_class
```

```
_warn_prf(average, modifier, msg_start, len(result))
```



335/335 [=====] - 1s 4ms/step - loss: 0.6943 - accuracy: 0.5055 - val\_loss: 0.6931 - val\_accuracy: 0.5034

Epoch 00003: val\_accuracy improved from 0.49664 to 0.50336, saving model to model\_save/weights-03-0.5034.hdf5  
val\_f1: 0.664238 - val\_precision: 0.497273 - val\_recall 1.000000 - val\_auc 0.500000  
Epoch 4/10

Epoch 00004: LearningRateScheduler setting learning rate to 0.08550000190734863.  
335/335 [=====] - 1s 4ms/step - loss: 0.6944 - accuracy: 0.4979 - val\_loss: 0.6931 - val\_accuracy: 0.5034


Epoch 00004: val\_accuracy did not improve from 0.50336

Epoch 00004: ReduceLROnPlateau reducing learning rate to 0.07695000171661377.  
val\_f1: 0.664238 - val\_precision: 0.497273 - val\_recall 1.000000 - val\_auc 0.500000  
Epoch 5/10

Epoch 00005: LearningRateScheduler setting learning rate to 0.07694999873638153.  
335/335 [=====] - 1s 4ms/step - loss: 0.6944 - accuracy: 0.4971 - val\_loss: 0.6961 - val\_accuracy: 0.4966

Epoch 00005: val\_accuracy did not improve from 0.50336

Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.06925499886274337.  
val\_f1: 0.000000 - val\_precision: 0.000000 - val\_recall 0.000000 - val\_auc 0.500000  
Epoch 00005: early stopping  
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/\_classification.py:1272: UndefinedMetricWarning: Precision is ill-defined and being set to  
\_warn\_prf(average, modifier, msg\_start, len(result))  
<keras.callbacks.History at 0x7fa011cd4a50>



!kill 198

/bin/bash: line 0: kill: (198) - No such process

%tensorboard --logdir logs/fit/model2/



☐ Show data download links

☐ Ignore outliers in chart scaling

 Tooltip sorting method: default

Smoothing



0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs


Write a regex to filter runs

☐ 20211009-070808/train

☐ 20211009-070808/validation

Filter tags (regular expressions supported)

epoch\_accuracy

epoch\_accuracy  
tag: epoch\_accuracy

run to download

epoch\_loss

## Model 3

```
input_layer = Input(shape=(2, ))
```

```
layer1 = Dense(784, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(input_layer)
layer2 = Dense(512, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer1)
layer3 = Dense(256, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer2)
layer4 = Dense(64, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer3)
layer5 = Dense(32, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer4)
```

```
output = Dense(1, activation='sigmoid', kernel_initializer=tf.keras.initializers.HeUniform)(layer5)
```

```
model3 = Model(inputs = input_layer, outputs = output)
optimizer = tf.keras.optimizers.SGD(learning_rate=0.1, momentum=0.9, nesterov=True)
```

```

model3.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

log_dir="logs/fit/model3/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
reduce_lr = ReduceLROnPlateau(monitor = 'val_accuracy', factor=0.9, patience=1, verbose=1, mode='auto')

terminate_tr = TerminateOnNaN()
earlystop = EarlyStopping(monitor='val_accuracy', patience=2, verbose=1, mode='max')

lrschedule = LearningRateScheduler(changedLearningRate, verbose=1)
model3.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test), validation_split=0.2, callbacks=[checkpoint, reduce_lr, terminate_tr,

Epoch 00003: LearningRateScheduler setting learning rate to 0.09500000141561031.
335/335 [=====] - 1s 4ms/step - loss: 0.6535 - accuracy: 0.6154 - val_loss: 0.7362 - val_accuracy: 0.5381

Epoch 00003: val_accuracy did not improve from 0.58843

Epoch 00003: ReduceLROnPlateau reducing learning rate to 0.08549999892711639.
val_f1: 0.679632 - val_precision: 0.520311 - val_recall 0.979586 - val_auc 0.543138
Epoch 4/10

Epoch 00004: LearningRateScheduler setting learning rate to 0.08550000190734863.
335/335 [=====] - 1s 4ms/step - loss: 0.6412 - accuracy: 0.6327 - val_loss: 0.6318 - val_accuracy: 0.6455

Epoch 00004: val_accuracy improved from 0.58843 to 0.64552, saving model to model_save/weights-04-0.6455.hdf5
val_f1: 0.665851 - val_precision: 0.668407 - val_recall 0.663315 - val_auc 0.668909
Epoch 5/10

Epoch 00005: LearningRateScheduler setting learning rate to 0.08550000190734863.
335/335 [=====] - 1s 4ms/step - loss: 0.6328 - accuracy: 0.6367 - val_loss: 0.6254 - val_accuracy: 0.6504

Epoch 00005: val_accuracy improved from 0.64552 to 0.65037, saving model to model_save/weights-05-0.6504.hdf5
val_f1: 0.688681 - val_precision: 0.653709 - val_recall 0.727605 - val_auc 0.673176
Epoch 6/10

Epoch 00006: LearningRateScheduler setting learning rate to 0.0812250018119812.
335/335 [=====] - 1s 4ms/step - loss: 0.6284 - accuracy: 0.6466 - val_loss: 0.6320 - val_accuracy: 0.6478

Epoch 00006: val_accuracy did not improve from 0.65037

Epoch 00006: ReduceLROnPlateau reducing learning rate to 0.07310250028967857.
val_f1: 0.618842 - val_precision: 0.715142 - val_recall 0.545399 - val_auc 0.665255
Epoch 7/10

Epoch 00007: LearningRateScheduler setting learning rate to 0.07310250401496887.

```

```
335/335 [=====] - 1s 4ms/step - loss: 0.6240 - accuracy: 0.6566 - val_loss: 0.6292 - val_accuracy: 0.6541

Epoch 00007: val_accuracy improved from 0.65037 to 0.65410, saving model to model_save/weights-07-0.6541.hdf5
val_f1: 0.663225 - val_precision: 0.681453 - val_recall 0.645948 - val_auc 0.673637
Epoch 8/10

Epoch 00008: LearningRateScheduler setting learning rate to 0.07310250401496887.
335/335 [=====] - 1s 4ms/step - loss: 0.6175 - accuracy: 0.6563 - val_loss: 0.6498 - val_accuracy: 0.5873

Epoch 00008: val_accuracy did not improve from 0.65410

Epoch 00008: ReduceLROnPlateau reducing learning rate to 0.06579225361347199.
val_f1: 0.430323 - val_precision: 0.802161 - val_recall 0.294028 - val_auc 0.611149
Epoch 9/10

Epoch 00009: LearningRateScheduler setting learning rate to 0.0625026423484087.
335/335 [=====] - 1s 4ms/step - loss: 0.6209 - accuracy: 0.6587 - val_loss: 0.6264 - val_accuracy: 0.6511

Epoch 00009: val_accuracy did not improve from 0.65410

Epoch 00009: ReduceLROnPlateau reducing learning rate to 0.05625238046050072.
val_f1: 0.643550 - val_precision: 0.698537 - val_recall 0.596587 - val_auc 0.670958
Epoch 00009: early stopping
<keras.callbacks.History at 0x7fa0e4de73d0>
```

```
!kill 197
```

```
/bin/bash: line 0: kill: (197) - No such process
```

```
%tensorboard --logdir logs/fit/model3/
```

Reusing TensorBoard on port 6008 (pid 380), started 0:01:31 ago. (Use '!kill 380' to kill it.)

TensorBoard

SCALARS

GRAPHS

DISTRIBUTIONS

HISTOGRAMS

TIME SERIES

INACTIVE

☐ Show data download links

☐ Ignore outliers in chart scaling

Tooltip sorting  
method: default ▼

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

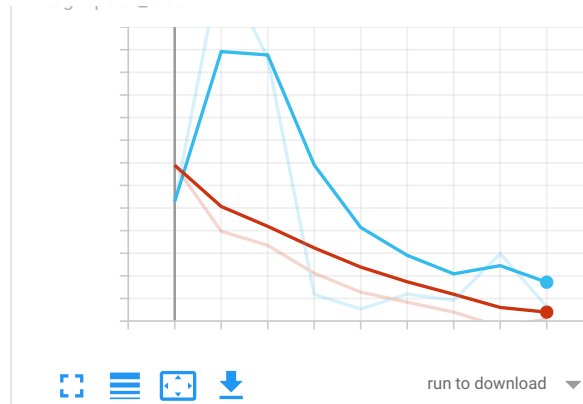
WALL

Runs

Write a regex to filter runs


☐ 20211009-070822/train

20211009-070822/train



epoch\_lr

epoch\_lr  
tag: epoch\_lr



## ▼ Model 4

20211009-070822/train

```
input_layer = Input(shape=(2, ))
```

```
layer1 = Dense(784, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(input_layer)
```

```
layer2 = Dense(512, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer1)
```

```
layer3 = Dense(256, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer2)
```

```
layer4 = Dense(64, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer3)
```

```
layer5 = Dense(32, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer4)
```

```
# layer6 = Dense(16, activation='relu', kernel_initializer=tf.keras.initializers.HeUniform)(layer5)
```

```
output = Dense(1, activation='sigmoid', kernel_initializer=tf.keras.initializers.HeUniform)(layer5)
```

```
model4 = Model(inputs=input_layer, outputs=output)
```

```

optimizer = tf.keras.optimizers.Adam(learning_rate=0.01, beta_1=0.9, beta_2=0.999, epsilon=1e-07, amsgrad=False, name="Adam",)

model4.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])

log_dir="logs/fit/model4/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=True, write_grads=True)

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')
reduce_lr = ReduceLROnPlateau(monitor='val_accuracy', factor=0.9, patience=1, verbose=1, mode='auto')

terminate_tr = TerminateOnNaN()
earlystop = EarlyStopping(monitor='val_accuracy', patience=2, verbose=1, mode='max')

lrschedule = LearningRateScheduler(changedLearningRate, verbose=1)
model4.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test), validation_split=0.2, callbacks=[checkpoint, reduce_lr, terminate_tr,

w 2.0 for the `TensorBoard` Callback.

.009999999776482582.
.4966 - accuracy: 0.5312WARNING:tensorflow:Callback method `on_train_batch_begin` is slow compared to the batch time (batch time: 0.0050s vs `on_train_batch_end` time: 0.0196s). Check your callbacks.
: 0.7274 - accuracy: 0.5904 - val_loss: 0.6610 - val_accuracy: 0.5955

ng model to model_save/weights-01-0.5955.hdf5
056 - val_auc 0.579095

.009999999776482582.
: 0.6340 - accuracy: 0.6368 - val_loss: 0.6334 - val_accuracy: 0.6466

aving model to model_save/weights-02-0.6466.hdf5
910 - val_auc 0.633827

.009499999787658453.
: 0.6326 - accuracy: 0.6509 - val_loss: 0.6178 - val_accuracy: 0.6638

aving model to model_save/weights-03-0.6638.hdf5
745 - val_auc 0.673258

.009499999694526196.
: 0.6119 - accuracy: 0.6605 - val_loss: 0.6070 - val_accuracy: 0.6728

aving model to model_save/weights-04-0.6728.hdf5
916 - val_auc 0.684422

```

```
.009499999694526196.  
: 0.6159 - accuracy: 0.6558 - val_loss: 0.6145 - val_accuracy: 0.6731
```

```
aving model to model_save/weights-05-0.6731.hdf5  
602 - val_auc 0.664810
```

```
.009024999709799886.  
: 0.6092 - accuracy: 0.6591 - val_loss: 0.6104 - val_accuracy: 0.6660
```

```
8122500032186508.  
553 - val_auc 0.674970
```

```
.008122500032186508.  
: 0.6084 - accuracy: 0.6582 - val_loss: 0.6056 - val_accuracy: 0.6724
```

```
73102500289678575.  
721 - val_auc 0.678912
```

```
!kill 197
```

```
/bin/bash: line 0: kill: (197) - No such process
```

```
%tensorboard --logdir logs/fit/model4/
```

Reusing TensorBoard on port 6006 (pid 378), started 0:01:07 ago. (Use '!kill 378' to kill it.)

## TensorBoard

SCALARS

GRAPHS

DISTRIBUTIONS

HISTOGRAMS

TIME SERIES

INACTIVE

- ☐ Show data download links
- ☐ Ignore outliers in chart scaling

Tooltip sorting  
method: default ▼

Smoothing

0.6

Horizontal Axis

STEP      RELATIVE      WALL


Runs

Write a regex to filter runs

- ☐ ○ 20211009-093418/train
- ☐ ○ 20211009-093418/validation
- ☐ ○ 20211009-093532/train
- ☐ ○ 20211009-093532/validation

TOGGLE ALL RUNS

logs/fit/model4/




run to download ▼

epoch\_loss



epoch\_loss  
tag: epoch\_loss



run to download ▼

epoch\_lr



epoch\_lr  
tag: epoch\_lr

